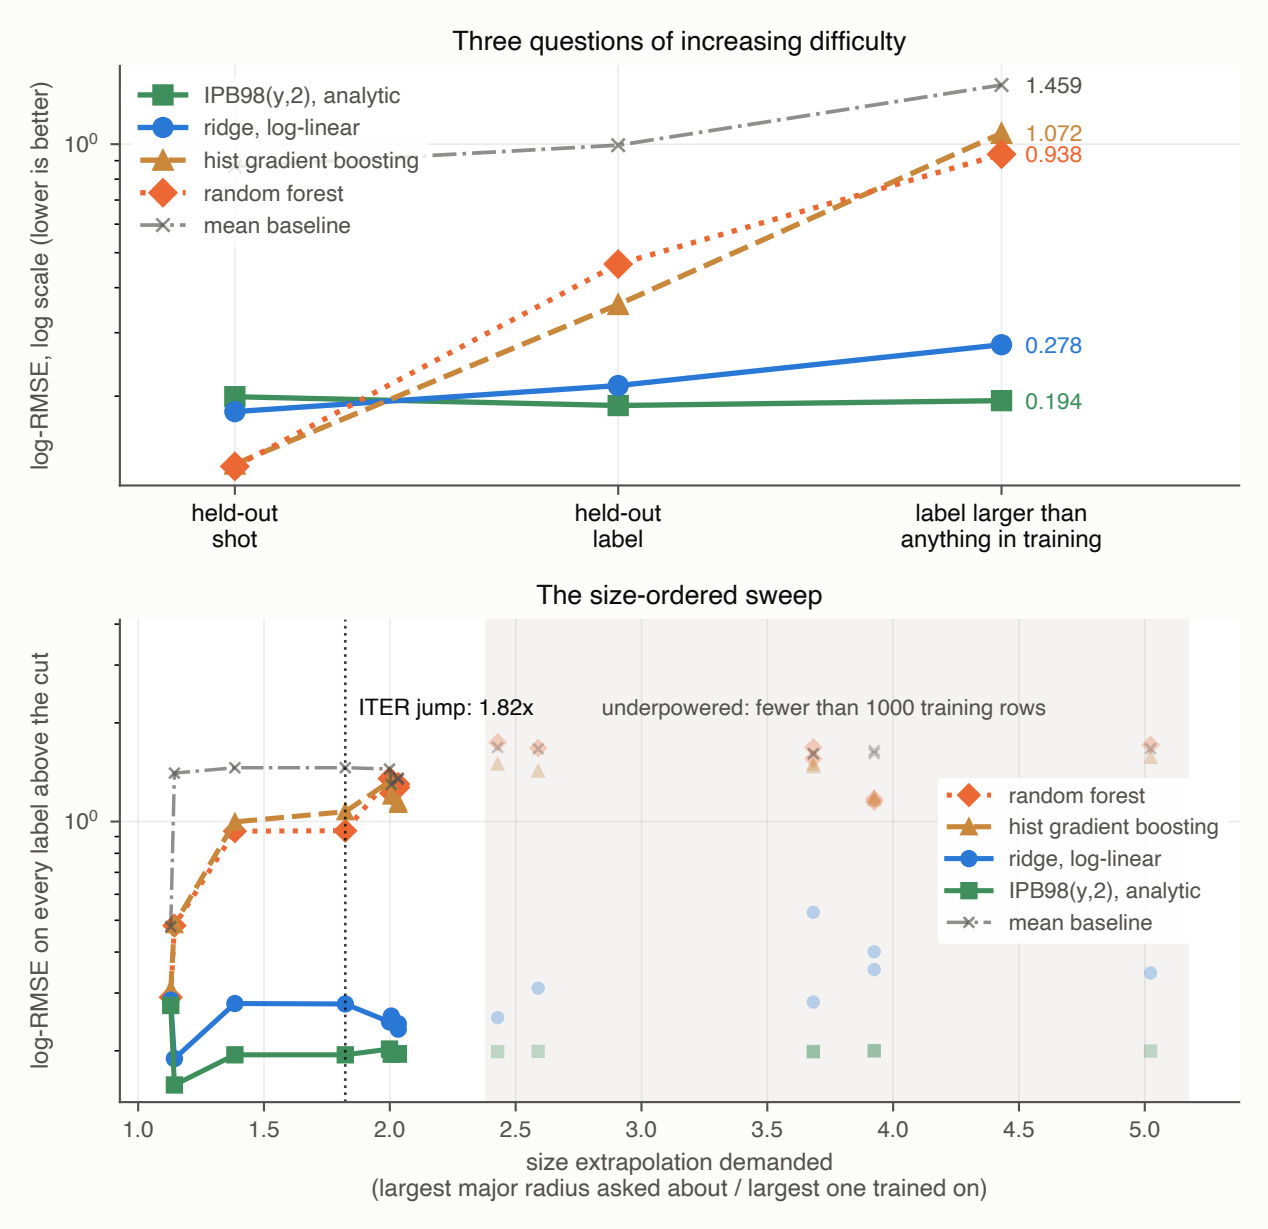
Three questions of increasing difficulty
log-RMSE, log scale (lower is better)
10^0
IPB98(y,2), analytic
ridge, log-linear
hist gradient boosting
random forest
mean baseline
1.459
1.072
0.938
0.278
0.194
held-out shot
held-out label
label larger than anything in training
The size-ordered sweep
log-RMSE on every label above the cut
10^0
ITER jump: 1.82x
underpowered: fewer than 1000 training rows
random forest
hist gradient boosting
ridge, log-linear
IPB98(y,2), analytic
mean baseline
1.0
1.5
2.0
2.5
3.0
3.5
4.0
4.5
5.0
size extrapolation demanded
(largest major radius asked about / largest one trained on)